

A Case for Relativistic Programming


Philip W. Howard

West Virginia University Institute of Technology

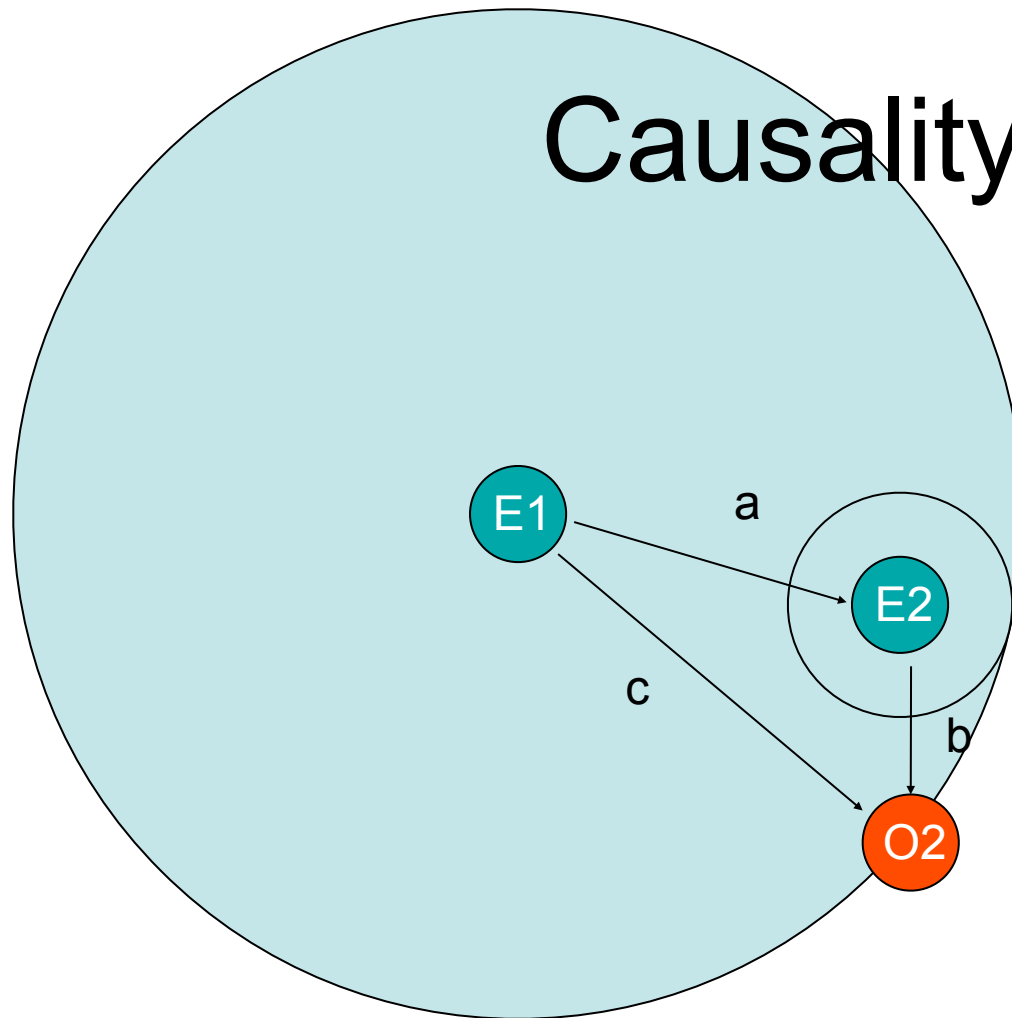
Jonathan Walpole

Portland State University

Events



Causality




Triangle inequality: $a + b \geq c$

Causality in Computers

- If computers were just wires, we'd be OK
- Computers are wires + illogic

illogic example




```
C = malloc();  
C->data = data;  
C->next = D;
```

```
B->next = C;
```

```
rp-publish(B->next, C);
```

Need for causal delays



```
B->next = D;  
wait-for-readers();  
free(C);
```

Rules for Placing Relativistic Programming Primitives

1. Write globally visible pointers with `rp_publish()`
2. When performing two writes, one earlier in traversal order and one later in traversal order, separate the two with `wait-for-readers()`

Relativistic vs. Concurrent Balanced Trees

- Relativistic Red Black Trees
- Concurrent AVL Trees

A practical concurrent binary search tree.
PPoPP '10; Bronson et al

Concurrent AVL get

```
20 V get(K k) {
21     return (V)attemptGet(k, rootHolder, 1, 0);
22 }
23
24 Object attemptGet(
25     K k, Node node, int dir, long nodeV) {
26     while (true) {
27         Node child = node.child(dir);
28         if (((node.version^nodeV) & IgnoreGrow) != 0)
29             return Retry;
30         if (child == null)
31             return null;
32         int nextD = k.compareTo(child.key);
33         if (nextD == 0)
34             return child.value;
35         long chV = child.version;
36         if ((chV & Shrinking) != 0) {
37             waitUntilNotChanging(child);
38         } else if (chV != Unlinked &&
39             child == node.child(dir)) {
40             if (((node.version^nodeV) & IgnoreGrow) != 0)
41                 return Retry;
42             Object p = attemptGet(k, child, nextD, chV);
43             if (p != Retry)
```

10/21/2012

Races'12

Relativistic RBTree Get


```
void *rb_find(rbtree *tree, long key)
{
    void *value;
    rbnode_t *node = tree->root;

    start_read();

    while (node != NULL) {
        if (key == node->key)
            break;
        else if (key < node->key)
            node = rp-dereference(node->left);
        else
            node = rp-dereference(node->right);
    }

    if (node != NULL)
        value = node->value;
    else
        value = NULL;
}
```

RBTree Delete



RP Delete

```
rbnode_t *new_node = rbnode_copy(swap);

rp-publish(new_node->left, node->left);
node->left->parent = new_node;

rp-publish(new_node->right, node->right);
node->right->parent = new_node;

if (is_left(node))
    rp-publish(prev->left, new_node);
else
    rp-publish (prev->right, new_node);
new_node->parent = prev;

// need to make sure new_node is seen before path to b is erased
wait-for-readers(tree->lock);

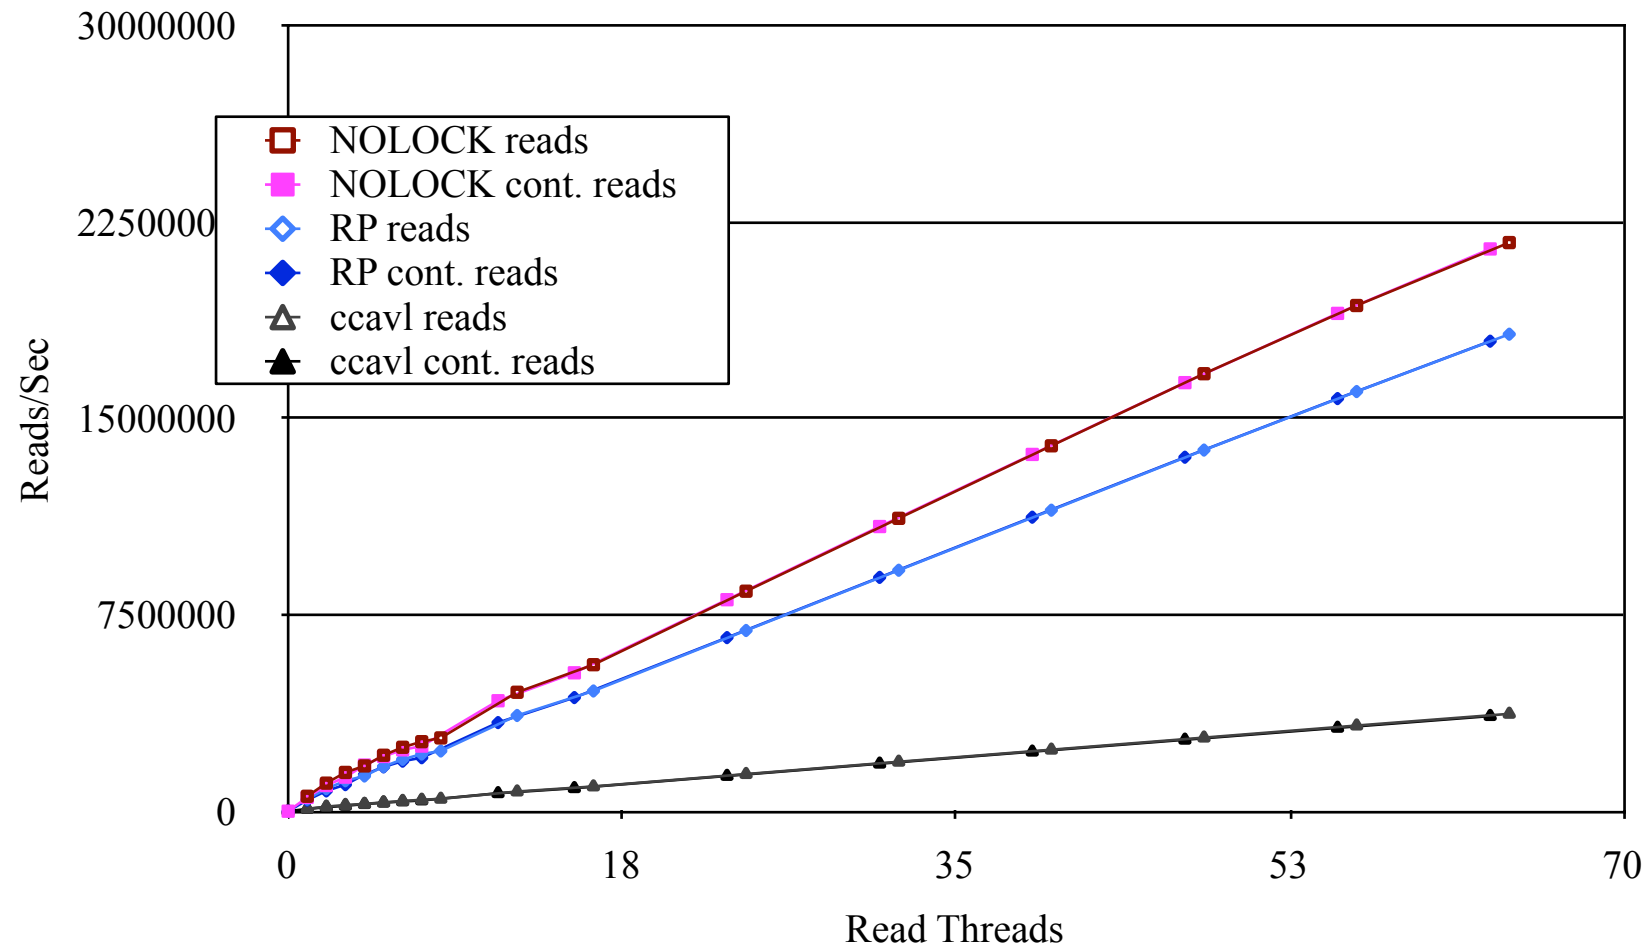
prev = swap->parent;
next = swap->right;

rp-publish(prev->left, swap->right);
if (swap->right != NULL) swap->right->parent = prev;
```

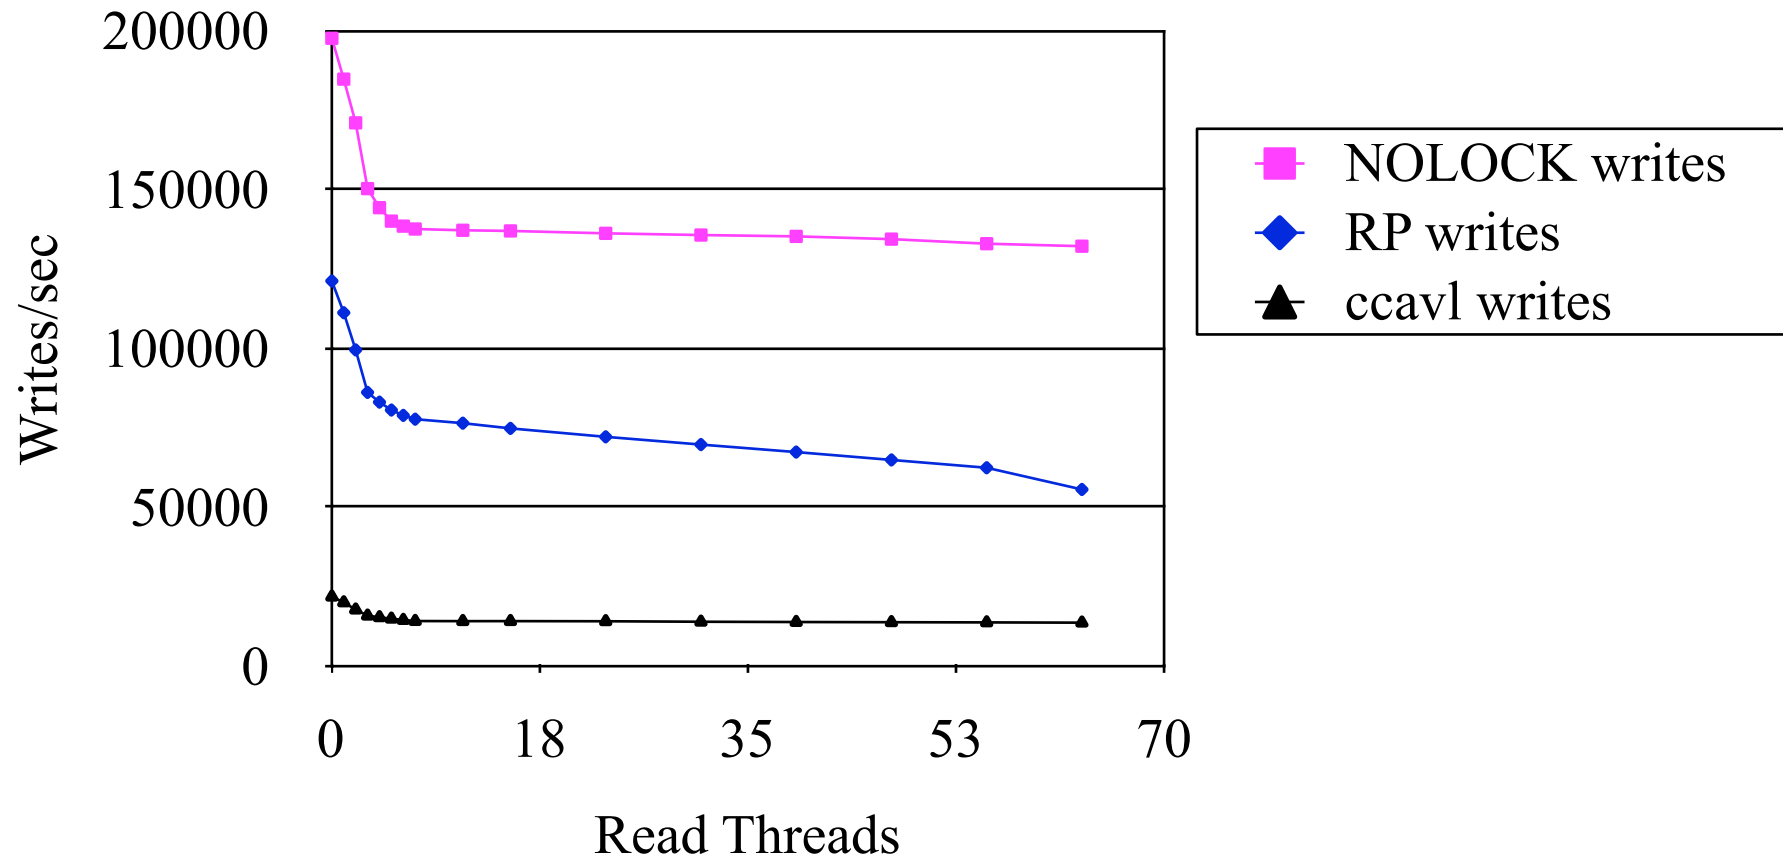
10/21/2012

Races'12

Read Performance



Contended Write Performance



Benefits of RP

- High performance, Highly scalable reads
- Simple Code
- Strong correctness properties