

Advanced Multiprocessor Programming

Practical lock implementations

Martin Wimmer

martin.wimmer@tuwien.ac.at

Research Group Parallel Computing

Faculty of Informatics, Institute of Information Systems

Vienna University of Technology (TU Wien)

Outline of this lecture

- General lock structure
- Bakery lock revisited
 - Common pitfalls in synchronization algorithms & motivating example
 - The unbounded timestamp problem
- Consensus (again)
- Locking algorithms based on higher consensus operations

Locks

- General structure
 - Lock operation
 - Unlock operation
 - Make sure unlock is always called even on errors
 - Java-style: in finally block
 - C++ style: RAII
- General distinction
 - Spinning/busy-waiting
 - We will concentrate on those
 - Blocking

```
// Java style
mutex.lock();
try {
    // critical section
} finally {
    mutex.unlock();
}
```

```
{ // C++ style
    std::lock_guard<std::mutex>
lock(mutex);
    // critical section
}
```

Bakery lock revisited

- As presented in previous lecture
 - Algorithm from Herlihy/Shavit book
 - Not the original Bakery algorithm
- Maintain first-come-first-served property
 - Idea: Number-dispensing machine
 - (like in bakeries)
 - Take a number. When your number is up, it is your turn.

```
// Initialize to false
bool flag[n];
// Initialize to 0
int label[n];

void lock() {
    int me = thread_id();
    flag[me] = true;
    // Acquire ticket
    label[me] =
        max(label[0], ... label[n - 1]);

    // Wait until no
    // lower ticket exists
    for(int i = 0; i < n; ++i) {
        while(i != me && flag[i] &&
            (label[i] < label[me] ||
             // On same label, take
             // thread with lower id
             (label[i] == label[me] &&
              i < me))) {
        }
    }
}

void unlock() {
    int me = thread_id();
    flag[me] = false;
}
```

Bakery lock mutual exclusion proof

- Proposition: The Bakery lock satisfies mutual exclusion
- Proof: By **contradiction** (from previous lecture)
 - labeling(A), labeling(B)
 - sequences of instructions generating the labels.
 - Assume $(\text{label}[A], A) << (\text{label}[B], B)$. When B entered it must therefore have read $\text{flag}[A] == \text{false}$

$\text{labeling}(B) \rightarrow \text{read}(B, \text{flag}[A] == \text{false}) \rightarrow \text{write}(A, \text{flag}[A] = \text{true}) \rightarrow \text{labeling}(A)$

- This contradicts $(\text{label}[A], A) << (\text{label}[B], B)$ since A's label would be at least $\text{label}[B] + 1$

Bakery lock proof

- Why do we revisit the proof again?

Bakery lock proof

- Why do we revisit the proof again?
- Because we forgot something
- The proof is not correct!
 - Fortunately the algorithm still seems to be correct after we correct the proof

Where is the error?

- Proposition: The Bakery lock satisfies mutual exclusion
- Proof: By **contradiction** (from previous lecture)
 - labeling(A), labeling(B)
 - sequences of instructions generating the labels.
 - Assume $(\text{label}[A], A) << (\text{label}[B], B)$. When B entered it must therefore have read $\text{flag}[A] == \text{false}$

$\text{labeling}(B) \rightarrow \text{read}(B, \text{flag}[A] == \text{false}) \rightarrow \text{write}(A, \text{flag}[A] = \text{true}) \rightarrow \text{labeling}(A)$

- This contradicts $(\text{label}[A], A) << (\text{label}[B], B)$ since A's label would be at least $\text{label}[B] + 1$

The error

- Labelling is not atomic!
 - Proof ignores this fact
 - Can be split into multiple reads and a write operation

```
flag[me] = true;

// read labels
int max = label[0];
for(int i = 1; i < n; ++i) {
    if(label[i] > max)
        max = label[i];
}

// write own label
label[me] = max + 1;

// Waiting part
...
```

- This is therefore not completely correct:
 - Assume $(\text{label}[A], A) << (\text{label}[B], B)$. When B entered it must therefore have read $\text{flag}[A] == \text{false}$
 $\text{labeling}(B) \rightarrow \text{read}(B, \text{flag}[A] == \text{false}) \rightarrow \text{write}(A, \text{flag}[A] = \text{true}) \rightarrow \text{labeling}(A)$
 - This path may also occur
 $\text{write}(B, \text{flag}[B] = \text{true}) \rightarrow \text{read}(B, \text{labels}) \rightarrow \text{write}(A, \text{flag}[A] = \text{true}) \rightarrow \text{read}(A, \text{labels}) \rightarrow \text{write}(B, \text{label}[B]) \rightarrow \text{read}(B, \text{flag}[A] == \text{true}) \rightarrow \text{read}(B, \text{label}[A]) \rightarrow \text{write}(A, \text{label}[A])$
 - This doesn't change correctness in this case
 - As $\text{label}[A]$ must be smaller than the new value, $\text{label}[A] < \text{label}[B]$ still holds
 - Still very dangerous to ignore such facts

Common pitfalls in synchronization algorithms

- Forgetting that some operation isn't atomic
 - As in the Bakery lock example
- Performing a check, and then assuming that the condition holds for the rest of the algorithm
- Overflow of timestamps
 - Next slide
- Assuming sequential consistency
 - We will come to this later
- Those problems are quite common
 - Might rarely lead to errors in practice
 - Random crashes that are very difficult to track
 - Trial and error is a bad idea => Make sure your algorithms are correct

Problem with unbounded registers

- Unbounded registers don't exist!
- Labels in Bakery algorithm may overflow, invalidating the first-come-first-served property
 - Common problem in synchronization algorithms
- Possible solutions
 - Ignore problem
 - make size of label large enough
 - _ e.g. on 64 bit systems
 - Simple hack:
 - _ use unsigned data-type
 - _ instead of checking $A > B$ check $A - B > 0$
 - _ Works well for many algorithms in practice, but may still fail in theory
 - _ Bad for Bakery, if some threads rarely access the lock
 - Construct unbounded timestamps
 - _ Cyclic graph providing an ordering of timestamps
 - _ May grow very large for larger numbers of threads
 - _ Quite complex to implement
 - Black-white bakery lock

Original Bakery lock

- Original Bakery algorithm
 - Flag field used to “lock” own label
- While thread is labelling, other threads wait for the updated value
 - This would allow us to treat labelling as atomic in the proof
(at least for this special case)
- Special value for label for unlocked state
 - To unlock we reset label
- We won't cover proof here

```
// Initialize to false
bool flag[n];
// Initialize to 0
int label[n];

void lock() {
    int me = thread_id();
    flag[me] = true;
    // Acquire ticket
    label[me] =
        max(label[0], ... label[n - 1]);
    flag[me] = false;

    // Wait until no
    // lower ticket exists
    for(int i = 0; i < n; ++i) {
        // Make sure thread is
        // finished choosing a label
        while(flag[i]);
        while(i != me &&
            label[i] != 0 &&
            (label[i] > label[me] ||
             // On same label, take
             // thread with lower id
             (label[i] == label[me] &&
              i < me)));
    }
}

void unlock() {
    int me = thread_id();
    label[me] = 0;
}
```

Black-white bakery lock

- Have a shared coloring bit
- Store color in addition to label
- Threads with color **different** to current bit come first
- After thread exits critical section, change shared coloring bit to color **different** to own color
- Guarantees first-come-first-served
- Only needs n values for label


```
void lock() {  
    ...  
    flag[me] = true;  
    mycolor[me] = color;  
    label[me] = ...  
    flag[me] = false;  
  
    for(int i = 0; i < n; ++i) {  
        // Make sure thread is  
        // finished choosing a label  
        while(flag[i]);  
        if(i != me) {  
            if(mycolor[me] == mycolor[i]) {  
                while([other wins] and  
                    mycolor[me] ==  
                    mycolor[i]);  
            } else {  
                while([other active] and  
                    mycolor[me] == color and  
                    mycolor[me] ==  
                    mycolor[i]);  
            }  
        }  
    }  
}  
  
void unlock() {  
    int me = thread_id();  
    color = !mycolor[me];  
    label[me] = 0;  
}
```

Waiting conditions

- Other thread has same color
 - Similar waiting conditions to original algorithm
 - If other thread changes color, we can also proceed
- Different color
 - Wait if my color equals shared color
 - Wait until
 - shared color changes
 - other thread unlocks
 - other thread changes colors

```
...
if(mycolor[me] == mycolor[i]) {
    // Same color case
    while(label[i] != 0 &&
        (label[i] < label[me] ||
        (label[i] == label[me] &&
        i < me)) &&
        mycolor[me] == mycolor[i]);
}
else {
    // Colors differ
    while(label[i] != 0 &&
        mycolor[me] == color &&
        mycolor[i] != mycolor[me]);
}
...
```

Black-white bakery lock



Still, the Bakery lock isn't practical

- $O(n)$ execution time
- $O(n)$ memory
- Can this be improved based on reading and writing memory?
 - No!
 - A lock for n threads needs at least n locations
(Recall proof in second lecture)
- Peterson and Bakery lock are thus optimal!
(but not practical)

Consensus

- Wait-free N-thread consensus protocol
 - Agree on a common value between n threads
 - Has to be wait free
 - This means the result may need to be decided even before some threads enter the protocol!
- Consensus number
 - A construction with consensus number n can be used to implement wait-free consensus for n threads

Consensus numbers

(Recall third lecture)

- Consensus number 1
 - Atomic registers
- Consensus number 2
 - Wait-free FIFO Queues
 - get-and-set
 - test-and-set
 - fetch-and-add
- Consensus number ∞
 - compare-and-swap (CAS)
 - load-linked + store conditional


Putting higher consensus operations to use

- Test-And-Set Lock
 - Single flag field
 - atomically set field to true if it isn't already
 - On success we have the lock
 - For unlock just reset the field
- Performance far from ideal
 - Each test-and-set call invalidates cached copies for all threads
 - High contention of interconnect

```
bool locked = false;

void lock() {
    // test_and_set checks whether
    // variable is set to false
    // and atomically sets it to
    // true
    while(!test_and_set(locked));
}

void unlock() {
    locked = false;
}
```




Test-And-Test-And-Set lock

- Only perform test-and-set if there is a chance of success
- Cache invalidated less often
- Still some contention with more threads

```
bool locked = false;

void lock() {
    while(true) {
        while(locked);
        if(test_and_set(locked))
            return true;
    }
}

void unlock() {
    locked = false;
}
```



Exponential Backoff

- On failure to acquire lock:
Backoff for a random amount of time
- Time to wait increases exponentially
- Reduces contention
 - On high contention threads try less often
 - Randomization makes sure threads wake up at different times
- Thread might wait longer than necessary!
- C++ note: Don't use the standard rand()
It uses locks inside!

```
class Backoff {  
    int limit = MIN_DELAY;  
    void backoff() {  
        int delay = rand() % limit;  
        limit = min(MAX_DELAY,  
                    limit*2);  
        sleep(delay);  
    }  
}
```

```
bool locked = false;  
  
void lock() {  
    Backoff bo;  
    while(true) {  
        while(locked);  
        if(test_and_set(locked))  
            return true;  
        bo.backoff();  
    }  
}  
  
void unlock() {  
    locked = false;  
}
```

Test-And-Set Lock summary

- Space complexity $O(1)$ for ∞ threads
 - Made possible by test-and-set
(consensus number 2)
- Problem with memory contention
 - All threads spin on a single memory location (cache coherence traffic)
- Threads might wait longer than necessary due to backoff
- Unfair/not starvation free

Queue lock


- First come – first served
- Less contention
 - Each thread spins on a local copy of a variable
 - (false sharing might occur, can be resolved with padding)
- Not space efficient!

```
bool flags[N] = {true, false,
                 false, false, ...}
int tail = 0;
thread_local int mySlot;

void lock() {
    mySlot =
        fetch_and_add(tail) % N;

    while(!flags[mySlot]) {};
}


void unlock() {
    flags[mySlot] = false;
    flags[(mySlot + 1) % N] = true;
}
```



CLH lock

- Linked list
- Single Sentinel node
- Spin on locked flag of previous node

```
Node* tail = new Node();  
thread_local Node* node;  
  
void lock() {  
    node = new Node();  
    node->locked = true;  
    node->pred =  
        get_and_set(&tail, node);  
    while(node->pred->locked) {}  
}  
  
void unlock() {  
    delete node->pred;  
    node->locked = false;  
}
```



CLH lock implementation notes

- C++ only supports static `thread_local` variables
 - Either pass on some data on lock and unlock
(if allowed by interface)
 - Or implement thread-local object storage yourself
- Differs slightly from Herlihy/Shavit
 - Predecessor stored in node
 - Manual memory management
 - We may safely delete predecessor after it is unlocked
(no other thread accessing it)

CLH lock properties

- First-come-first-served
- $O(L)$ space, where L is the number of threads currently accessing the lock
 - Some more space depending on implementation of `thread_local` data
 - Herlihy/Shavit implementation requires $O(p)$
- Each thread spins on a separate location
 - Allocated locally by each thread, reduces false sharing
- Problem on some architectures:
 - locked field is in remote location

MCS lock

- Reverse list
- To unlock, modify locked field of next node
- If no successor exists, reset tail

```
Node* tail = nullptr;
thread_local Node* node =
    new Node();

void lock() {
    Node* n = node;
    pred = get_and_set(&tail, n);
    if(pred != nullptr) {
        n->locked = true;
        pred->next = n;
        while(n->locked);
    }
}

void unlock() {
    Node* n = node;
    if(n->next == nullptr) {
        if(CAS(&tail, n, nullptr))
            return;
        // Wait for next thread
        while(n->next == null);
    }
    n->next->locked = false;
    n->next = nullptr;
}
```



MCS lock properties

- First-come-first-served
- $O(p)$ space
 - Some more space depending on implementation of `thread_local` data
- Each thread spins on its own memory location
 - Updated by other thread
- Requires compare-and-swap
 - (consensus number ∞)
- Unlock is not wait-free any more!
 - We might have to wait for the next lock owner to set the next pointer

Locks with timeouts

- Abandoning is easy for Test-And-Set lock
 - Just stop trying to acquire lock
 - Timing out is wait-free
- More difficult for queue locks
 - If we just exit, the following thread will starve
 - Can't just unlink the node
(other thread might be accessing it)
- Lazy approach
 - Mark node as abandoned
 - Successor is responsible for cleanup

```
Node* tail = nullptr;
thread_local Node* node;

void lock() {
    node = new Node();
    node->locked = true;
    node->pred =
        get_and_set(&tail, node);
    while(node->pred != nullptr &&
        node->pred->locked) {

        // Remove abandoned nodes
        if(node->pred->abandoned) {
            Node* tmp =
                node->pred->pred;
            delete node->pred;
            node->pred = tmp;
        }
    }
}

void timeout() {
    node->abandoned = true;
}

void unlock() {
    if(node->pred != nullptr)
        delete node->pred;
    node->locked = false;
}
```

Composite lock

- Combines Backoff lock and Queue lock
- Preallocate fixed number of nodes $< p$
- Acquire a random node
 - Only one thread may use a certain node
 - On failure back off
- As soon as a node is acquired, enqueue it
- Can be augmented with a fast path for low contention
 - If queue is empty, try fast path, on failure use normal path

```
Node nodes[n];

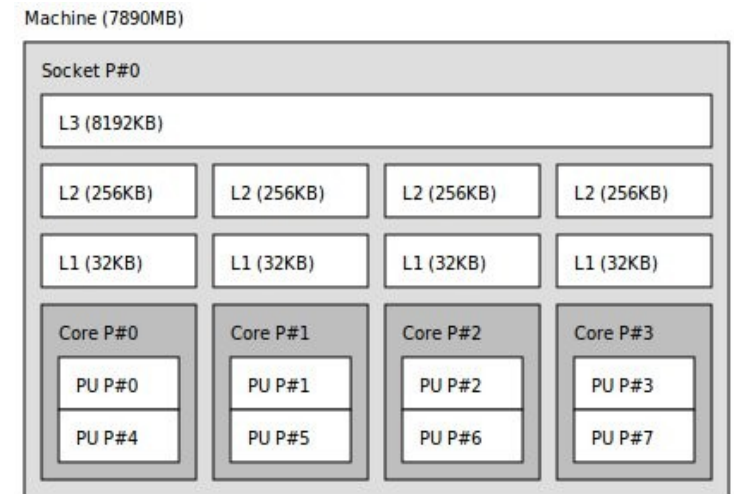
void lock() {
    Backoff bo;
    Node* node;
    while(!(n =
        acquireNode(rand() % n)));

    enqueueNode(n);

    wait(n);
}
```

Hierarchical locks


- Most modern architectures can be modeled in a hierarchy
 - Some processors are “near” to each other
 - Smaller memory access times



- We can improve efficiency by taking advantage of this
- Look at the architecture of your own machine
 - Istopo tool (part of hwloc)
 - Try it on saturn as well


Hierarchical backoff lock

- Make distinction between local and remote accesses
- Use shorter backoff for local accesses
 - Local threads are more likely to acquire lock
- May starve remote threads
- There may be more than two levels in the hierarchy!



Hierarchical queue lock

- Build queues locally
- In addition a global queue exists
- First thread in local queue is cluster master
 - tries to splice in local queue into global queue
- First element added to local queue after splicing in is new cluster master



What we have learned so far

- To implement efficient locks, higher consensus operations are needed
 - At least consensus number 2
- Difficulties
 - Fairness
 - Congestion
 - Memory hierarchy
 - Space considerations
 - Integer overflow

What we ignored so far

- Memory consistency (coming next)
 - We assumed sequential consistency
 - Sequential consistency is not very efficient on current architectures
 - Programming with other consistency models is more difficult

Project topic

- Implement various locks
 - Test-And-Set lock
 - Test-And-Test-And-Set lock
 - Backoff lock
 - Black-White Bakery lock
 - _ Try different consistency models of C++11 (sequential consistency, acquire/release)
 - _ Also implement with volatile keyword and fences for comparison
 - CLH lock
 - Hierarchical Backoff lock
 - _ For n levels
 - _ distance function provided by framework
- Implementation as single classes in C++
 - Should conform to a certain interface (tbd – C++11 compatible)
 - We will provide test program in Pheet framework where locks can be plugged in
 - Run performance measurements with test program on saturn
- Report experiences & analyze performance
 - Short written report (what you learned, difficulties, pitfalls, performance analysis) + plots
 - Probably short presentation at end of semester (share experience with other groups)